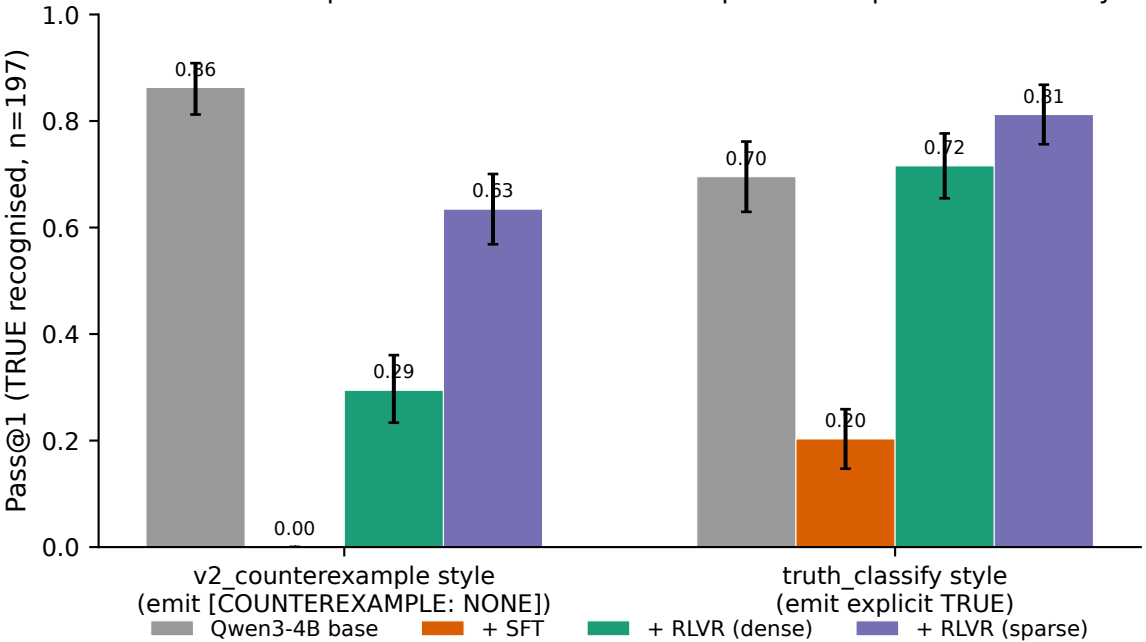
Theorem-truth probe — the SFT imitation trap and the sparse-RL recovery
Pass@1 (TRUE recognised, n=197)
1.0
0.8
0.6
0.4
0.2
0.0
0.86
0.00
0.29
0.63
0.70
0.20
0.72
0.81
v2_counterexample style
(emit [COUNTEREXAMPLE: NONE])
truth_classify style
(emit explicit TRUE)
Qwen3-4B base
+ SFT
+ RLVR (dense)
+ RLVR (sparse)